

EXPERIMENT No: 5

Title : Study & Implementation of

- Group by & Having Clause
- Order by Clause
- Indexing
- Sub queries
- Views

Theory:

GROUP BY:

The GROUP BY clause groups rows that have the same values into summary rows.

Syntax: SELECT column1, aggregate_function(column2)

FROM table_name

GROUP BY column1;

Example 1: Simple Grouping

– Count employees in each department

SELECT department_id, COUNT(*) as employee_count

FROM employees

GROUP BY department_id;

Example 2:

– Calculate the sum of salaries of all employees

SELECT EMPNO, SUM (SALARY) as Total_Salary

FROM EMPLOYEES

GROUP BY EMPNO;

Example 3: Grouping by Multiple Columns

– Count employees by department and job title

SELECT department_id, job_id, COUNT(*) as employee_count

FROM employees

GROUP BY department_id, job_id;

HAVING :

The HAVING clause was added to SQL because the WHERE keyword could not be used with aggregate functions. The HAVING clause must follow the GROUP BY clause in a query and must also precede the ORDER BY clause if used.

Syntax: SELECT column1, aggregate_function(column2)

FROM table_name

GROUP BY column1

HAVING condition;

Example 1: Filtering Groups

– Find departments with more than 5 employees

SELECT department_id, COUNT(*) as employee_count

FROM employees

GROUP BY department_id

HAVING COUNT(*) > 5;

Example 2: HAVING with Multiple Conditions

– Find departments with avg salary > 5000 and at least 3 employees

SELECT department_id, AVG(salary) as avg_salary, COUNT(*) as emp_count

FROM employees



Edit with WPS Office

```
GROUP BY department_id  
HAVING AVG(salary) > 5000 AND COUNT(*) >= 3;
```

ORDER BY:

The ORDER BY clause sorts the result set in ascending or descending order.

Syntax:

```
SELECT column1, column2  
FROM table_name  
ORDER BY column1 [ASC|DESC], column2 [ASC|DESC];
```

Example 1: Simple Sorting

– Sort employees by last name

```
SELECT employee_id, last_name, first_name  
FROM employees  
ORDER BY last_name;
```

Example 2: Sorting by Multiple Columns

– Sort by department then by salary (highest first)

```
SELECT employee_id, last_name, department_id, salary  
FROM employees  
ORDER BY department_id, salary DESC;
```

Combining GROUP BY, HAVING, and ORDER BY

– Find departments with total salary > 20000, order by total salary descending

```
SELECT department_id, SUM(salary) as total_salary, COUNT(*) as  
employee_count  
FROM employees  
GROUP BY department_id  
HAVING SUM(salary) > 20000  
ORDER BY total_salary DESC;
```

– Find job titles with average salary between 5000 and 10000, ordered by avg salary

```
SELECT job_id, AVG(salary) as avg_salary, COUNT(*) as employee_count  
FROM employees  
GROUP BY job_id  
HAVING AVG(salary) BETWEEN 5000 AND 10000  
ORDER BY avg_salary DESC;
```

INDEXING:

Indexes in Oracle are database objects that improve data retrieval performance. They work like a book's index, allowing the database to find data without scanning the entire table.

Types of Indexes in Oracle

1) **B-Tree Index(Balanced Tree Index):** The **most common type** of index, suitable for high-cardinality columns (columns with many distinct values).

Example:

– Create a basic B-Tree index

```
CREATE INDEX emp_last_name_idx ON employees(last_name);
```



Edit with WPS Office

– **Composite index on multiple columns**

```
CREATE INDEX emp_dept_job_idx ON employees(department_id, job_id);
```

2) **Bitmap Index:** Best for low-cardinality columns (columns with few distinct values like gender or status flags).

– **Create a bitmap index on gender column**

```
CREATE BITMAP INDEX emp_gender_idx ON employees(gender);
```

– **Bitmap index on department_id (if few departments)**

```
CREATE BITMAP INDEX emp_dept_bitmap_idx ON employees(department_id);
```

3) **Function-Based Index:** Indexes the result of a function or expression.

```
CREATE INDEX idx_employee_upper_name  
ON employees (UPPER(last_name));
```

4) **Unique Index:** Ensures that all values in the indexed column are unique.

```
CREATE UNIQUE INDEX idx_employee_email  
ON employees (email);
```

5) **Composite Index:** Index on multiple columns.

```
CREATE INDEX idx_employee_name_dept  
ON employees (last_name, department_id);
```

6) **Reverse Key Index:** Reverses the bytes of the indexed column, useful for balancing I/O in certain scenarios.

```
CREATE INDEX idx_employee_id_reverse  
ON employees (employee_id) REVERSE;
```

Index Management Examples

View existing indexes:

```
SELECT index_name, index_type, table_name  
FROM user_indexes  
WHERE table_name = 'EMPLOYEES';
```

Dropping an Index:

To remove an index, use the DROP INDEX statement:

```
DROP INDEX index_name;  
DROP INDEX emp_last_name_idx;
```

Rebuild an index (for maintenance):

```
ALTER INDEX emp_last_name_idx REBUILD;
```

Gather statistics for an index:

```
ANALYZE INDEX emp_last_name_idx COMPUTE STATISTICS;
```

Rename an Index

```
ALTER INDEX idx_employee_name RENAME TO idx_emp_name;
```

When to Use Indexes

Columns frequently used in WHERE clauses.

Columns used in JOIN conditions.

Columns with high selectivity (many distinct values).

When Not to Use Indexes

On small tables.

On columns with low selectivity (few distinct values).

On columns frequently updated (indexes slow down INSERT, UPDATE, and DELETE operations).

Performance Considerations



Edit with WPS Office

Example: Force index usage (hint):

```
SELECT /*+ INDEX(employees emp_last_name_idx) */ *  
FROM employees  
WHERE last_name = 'Smith';
```

Example: Check if index is used (explain plan):

```
EXPLAIN PLAN FOR  
SELECT * FROM employees WHERE last_name = 'Smith';  
  
SELECT * FROM TABLE(DBMS_XPLAN.DISPLAY);
```

Partial/Filtered index (Oracle 12c+):

```
CREATE INDEX emp_active_idx ON employees(employee_id)  
WHERE status = 'ACTIVE';
```

SUBQUERIES:

Subqueries (or nested queries) are queries embedded within other SQL statements. They are powerful tools for complex data retrieval in Oracle.

Types of Subqueries

1. Single Row Subquery

Returns a single row and is used with single-row comparison operators like =, >, <, etc.

Find employees who earn more than the average salary.

```
SELECT employee_id, first_name, last_name, salary  
FROM employees  
WHERE salary > (SELECT AVG(salary) FROM employees);
```

2. Multi-Row Subqueries

Returns multiple rows

– Find employees who work in the same department as employees with the last name 'Smith'.

```
SELECT employee_id, first_name, last_name, department_id  
FROM employees  
WHERE department_id IN (SELECT department_id FROM employees WHERE  
last_name = 'Smith');
```

– Find employees in departments located in the US

```
SELECT employee_id, last_name, department_id  
FROM employees  
WHERE department_id IN (  
                SELECT department_id FROM departments  
WHERE location_id IN (  
                SELECT location_id FROM locations  
WHERE country_id = 'US' ));
```

3 .Multi-column Subqueries

Return multiple columns.

– Find employees with same job and department as employee 103

```
SELECT employee_id, last_name, job_id, department_id  
FROM employees  
WHERE (job_id, department_id) = (  
        SELECT job_id, department_id  
        FROM employees  
        WHERE employee_id = 103);
```



```
WHERE employee_id = 103  
);
```

4. Correlated subquery

A subquery that references a column from the outer query.

Find employees who earn more than the average salary in their department.

```
SELECT employee_id, first_name, last_name, salary, department_id  
FROM employees e1  
WHERE salary > (SELECT AVG(salary) FROM employees e2 WHERE  
e2.department_id = e1.department_id);
```

5. Scalar Subquery: Returns a single value and can be used in the SELECT clause or as part of an expression.

Retrieve the employee's name and the average salary of their department.

```
SELECT first_name, last_name, salary,  
       (SELECT AVG(salary) FROM employees e2 WHERE e2.department_id =  
e1.department_id) AS avg_dept_salary  
FROM employees e1;
```

VIEW:

Views are virtual tables that represent the result of a stored query. They don't store data physically but provide a way to simplify complex queries, enhance security, and present data differently to different users.

A view is a virtual table, which consists of a set of columns from one or more tables. It is similar to a table but it does not store in the database. View is a query stored as an object.

Syntax: CREATE VIEW <view_name> AS SELECT <set of fields>
FROM relation_name WHERE (Condition)

Example:

```
SQL> CREATE VIEW employee AS SELECT empno,ename,job FROM EMP  
WHERE job = 'clerk';  
SQL> View created.
```

Example:

```
CREATE VIEW [Current Product List] AS  
SELECT ProductID, ProductName  
FROM Products  
WHERE Discontinued=No;
```

UPDATING A VIEW : A view can updated by using the following syntax :

Syntax : CREATE OR REPLACE VIEW view_name AS
SELECT column_name(s)
FROM table_name
WHERE condition

DROPPING A VIEW: A view can deleted with the DROP VIEW command.

Syntax: DROP VIEW <view_name> ;

LAB PRACTICE ASSIGNMENT:

Consider the following schema:



Edit with WPS Office

Sailors (sid, sname, rating, age)

Boats (bid, bname, color)

Reserves (sid, bid, day(date))

1. Creating the Tables

-- Creating Sailors Table

```
CREATE TABLE Sailors
(
    sid NUMBER(5) PRIMARY KEY,
    sname VARCHAR2(50) NOT NULL,
    rating NUMBER(2) CHECK (rating >= 1 AND rating <= 10),
    age NUMBER(3,1) CHECK (age > 0)
);
```

-- Creating Boats Table

```
CREATE TABLE Boats
(
    bid NUMBER(5) PRIMARY KEY,
    bname VARCHAR2(50) NOT NULL,
    color VARCHAR2(20) NOT NULL
);
```

-- Creating Reserves Table

```
CREATE TABLE Reserves
(
    sid NUMBER(5),
    bid NUMBER(5),
    day DATE,
    PRIMARY KEY (sid, bid, day),
    FOREIGN KEY (sid) REFERENCES Sailors(sid) ON DELETE CASCADE,
    FOREIGN KEY (bid) REFERENCES Boats(bid) ON DELETE CASCADE
);
```

2. Inserting Sample Records

-- Inserting data into Sailors table

```
INSERT ALL
    INTO Sailors VALUES (1, 'Alice', 7, 23.5)
    INTO Sailors VALUES (2, 'Bob', 5, 29.0)
    INTO Sailors VALUES (3, 'Charlie', 6, 35.2)
    INTO Sailors VALUES (4, 'David', 8, 40.0)
    INTO Sailors VALUES (5, 'Eve', 3, 21.7)
SELECT * FROM DUAL;
```

-- Inserting data into Boats table

```
INSERT ALL
    INTO Boats VALUES (101, 'Seafarer', 'red')
    INTO Boats VALUES (102, 'Ocean Wave', 'blue')
    INTO Boats VALUES (103, 'Sunshine', 'green')
    INTO Boats VALUES (104, 'Speedster', 'red')
    INTO Boats VALUES (105, 'Blue Pearl', 'blue')
SELECT * FROM DUAL;
```

-- Inserting data into Reserves table

```
INSERT ALL
    INTO Reserves VALUES (1, 101, TO_DATE('2025-03-01', 'YYYY-MM-DD'))
    INTO Reserves VALUES (2, 103, TO_DATE('2025-03-02', 'YYYY-MM-DD'))
```



Edit with WPS Office

```

INTO Reserves VALUES (3, 104, TO_DATE('2025-03-01', 'YYYY-MM-DD'))
INTO Reserves VALUES (4, 102, TO_DATE('2025-03-03', 'YYYY-MM-DD'))
INTO Reserves VALUES (5, 101, TO_DATE('2025-03-04', 'YYYY-MM-DD'))
INTO Reserves VALUES (2, 105, TO_DATE('2025-03-05', 'YYYY-MM-DD'))
INTO Reserves VALUES (3, 101, TO_DATE('2025-03-06', 'YYYY-MM-DD'))
INTO Reserves VALUES (4, 103, TO_DATE('2025-03-06', 'YYYY-MM-DD'))
SELECT * FROM DUAL;

```

1. Find all information of sailors who have reserved boat number 101.

```

SELECT S.* FROM Sailors S
JOIN Reserves R ON S.sid = R.sid
WHERE R.bid = 101;

```

sid	sname	rating	age
1	Alice	7	23.5
5	Eve	3	21.7
3	Charlie	6	35.2

2. Find the name of boat reserved by Bob.

```

SELECT DISTINCT B.bname FROM Boats B
JOIN Reserves R ON B.bid = R.bid
JOIN Sailors S ON S.sid = R.sid
WHERE S.sname = 'Bob';

```

bname
Sunshine
Blue Pearl

3. Find the names of sailors who have reserved a red boat, and list in the order of age.

```

SELECT S.sname,S.age FROM Sailors S
JOIN Reserves R ON S.sid = R.sid
JOIN Boats B ON R.bid = B.bid
WHERE B.color = 'red'
ORDER BY S.age;

```

sname
Eve
Alice
Charlie

4. Find the names of sailors who have reserved at least one boat.

```

SELECT DISTINCT S.sname FROM Sailors S
JOIN Reserves R ON S.sid = R.sid;

```

sname
Alice
Bob
Charlie
David
Eve

5. Find the ids of sailors who have reserved a red boat or a green boat.

```

SELECT DISTINCT R.sid FROM Reserves R
JOIN Boats B ON R.bid = B.bid

```

WHERE B.color IN ('red', 'green');

sid
1
2
3
4

6. Find the name and the age of the youngest sailor.

SELECT sname, age FROM Sailors
WHERE age = (SELECT MIN(age) FROM Sailors);

sname	age
Eve	21.7

7. Count the number of different sailor names

SELECT COUNT(DISTINCT sname) AS unique_sailor_names
FROM Sailors;

unique_sailor_names	
5	

8. Find the average age of sailors for each rating level.

SELECT rating, AVG(age) AS avg_age FROM Sailors
GROUP BY rating;

rating	avg_age
3	21.7
5	29.0
6	35.2
7	23.5
8	40.0

9. Find the average age of sailors for each rating level that has at least two sailors.

SELECT rating, AVG(age) AS avg_age
FROM Sailors
GROUP BY rating
HAVING COUNT(sid) >= 2;

rating	avg_age
5	28.5
6	34.7

